

shredCVL using Logistic regression in R

1 March, 2018

Contents

Introduction	1
Exploratory data analysis	1
Feature engineering	12
Logistic Regression	15
Training and testing sets	15
Training logistic regressions	15

Introduction

*****For the second project we'll explore user data from shredCVL to identify accounts likely belonging to bots. The data set has variables about profile configuration (default_profile, default_profile_image), connectivity (friendsCount, followersCount), and some information about the nature of their tweets (diversity, mean_mins_between_tweets). Additionally, there's an outcome variable called bot that denotes whether the account belongs to a bot (bot == 1) or to a human (bot == 0).*****

Exploratory data analysis

We've got a brand new data set, so let's familiarize ourselves by conducting an exploratory data analysis. Let's start by summarizing the whole data set to see what the variable values are.

```
shredCVL <- readRDS("capstone dataset/shredCVL")
```

```
summary(shredCVL)
```

```
##   screenName          bot      statusesCount    friendsCount
## Length: 2618      Min.   : 0.000      Min.   :      0      Min.   :      0
## Class : character  1st Qu.: 0.000      1st Qu.:   2524     1st Qu.:   258
## Mode  : character  Median : 0.000      Median :   13042    Median :    661
##                               Mean   : 0.195      Mean   :   45605    Mean   :   1683
##                               3rd Qu.: 0.000      3rd Qu.:   41709    3rd Qu.:   1678
##                               Max.   : 1.000      Max.   :  1742257    Max.   :   68232
##                               NA s   : 29         NA s   : 29
## followersCount     listedCount      acct_created
## Min.   :      0      Min.   :      0      Min.   : 2007-02-08 04:24:56
## 1st Qu.:    277     1st Qu.:      4     1st Qu.: 2010-04-16 14:56:21
## Median :    915     Median :     22     Median : 2012-09-06 16:14:22
## Mean   :   10366     Mean   :    143     Mean   : 2012-11-08 23:14:03
## 3rd Qu.:   2296     3rd Qu.:    104     3rd Qu.: 2015-07-21 08:54:27
## Max.   :  6604309     Max.   :   39981     Max.   : 2017-12-11 13:58:41
## NA s   : 29         NA s   : 29
## julianCreated      acct_age      langDiv
## Length: 2618      Min.   :   1361      Min.   : 0.39
```

```
## Class : difftime      1st Qu.: 766986      1st Qu.: 2.95
## Mode : numeric      Median : 2275666      Median : 5.57
##                               Mean : 2221613      Mean : 5.59
##                               3rd Qu.: 3534304      3rd Qu.: 7.50
##                               Max. : 5209655      Max. : 29.07
##                               NA s : 193
## mean_time_betwn_tweets      App      Count      App. BoN
## Min. : 0      Length: 2618      Min. : 1      Min. : 0.00
## 1st Qu.: 40      Class : character      1st Qu.: 62      1st Qu.: 0.00
## Median : 127      Mode : character      Median : 412      Median : 0.00
## Mean : 12156                               Mean : 772      Mean : 0.03
## 3rd Qu.: 549                               3rd Qu.: 657      3rd Qu.: 0.00
## Max. : 3598898                               Max. : 6238      Max. : 1.00
## NA s : 38      NA s : 164      NA s : 164
## mCount
## Min. : 1.00
## 1st Qu.: 1.00
## Median : 1.00
## Mean : 1.29
## 3rd Qu.: 1.00
## Max. : 4.00
##
```

From the summary, we can see that there are a couple factor variables in the data set, bot, default_t_profile, default_t_profile_image and geo_enabled. Before exploring further, let's first tell R that those columns represent categorical variables.

```
names(shredCVL)
```

```
## [1] "screenName"      "bot"
## [3] "statusesCount"   "friendsCount"
## [5] "followersCount"  "listedCount"
## [7] "acct_created"    "julianCreated"
## [9] "acct_age"        "langDiv"
## [11] "mean_time_betwn_tweets" "App"
## [13] "Count"           "App. BoN"
## [15] "mCount"
```

```
shredCVL$bot = factor(shredCVL$bot)
shredCVL$App.BoN = factor(shredCVL$App.BoN)
```

```
shredCVL <- shredCVL %>%
```

```
  select(statusesCount, friendsCount, followersCount, listedCount, acct_age, langDiv, mean_time_betwn_tweets,
```

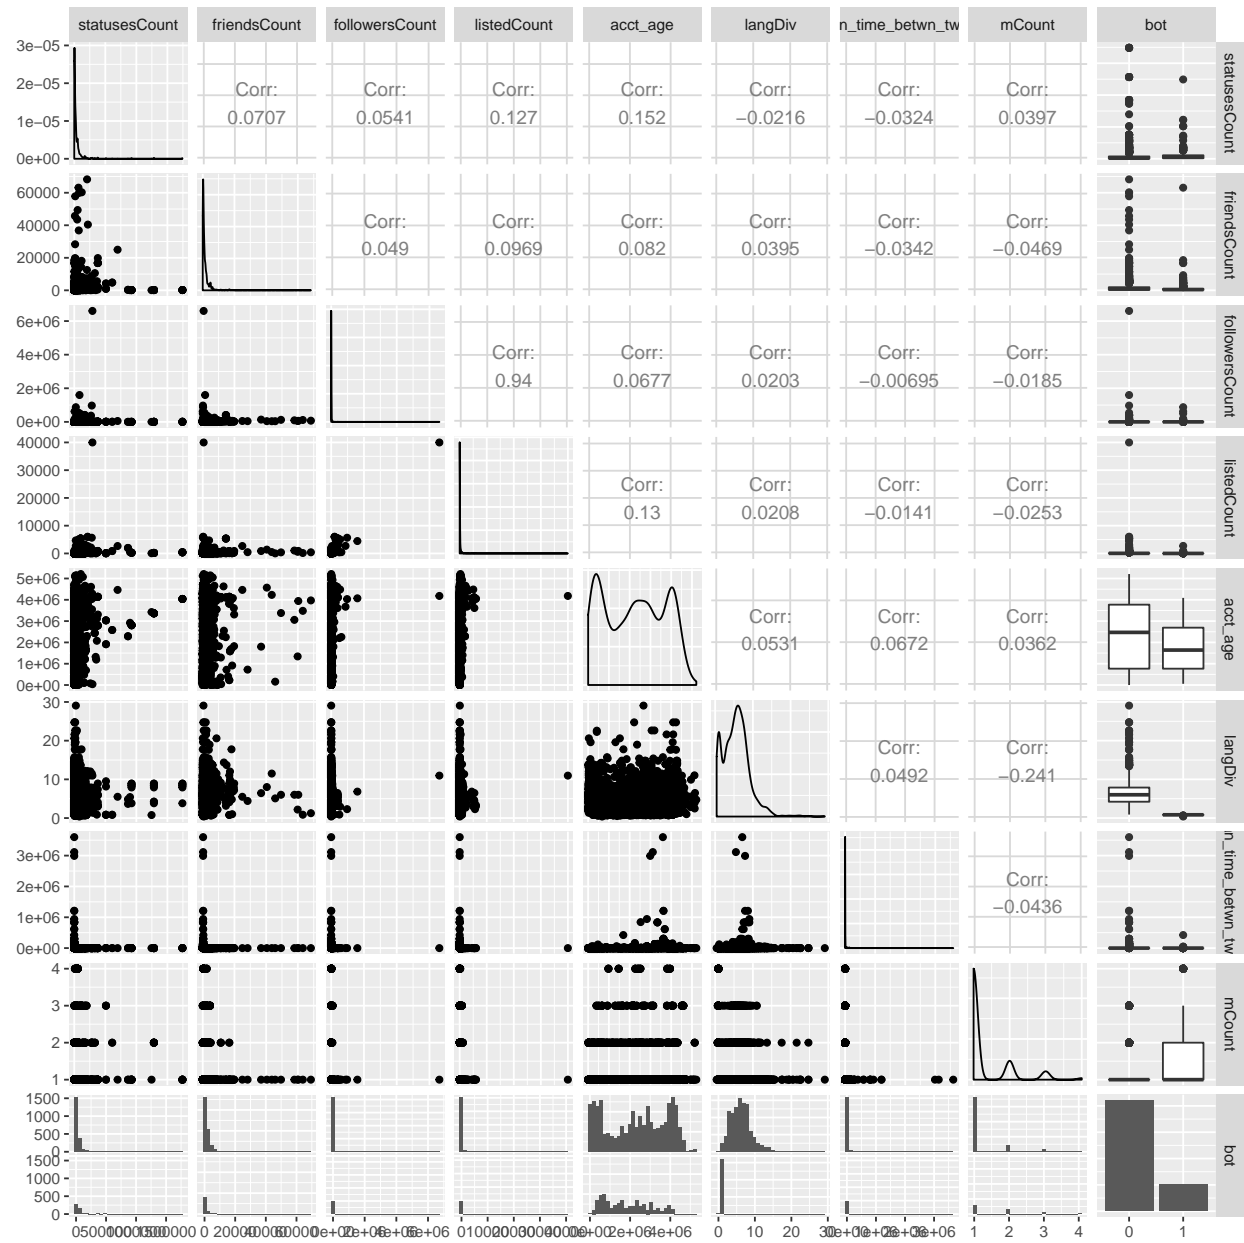
```
summary(shredCVL)
```

```
## statusesCount      friendsCount      followersCount      listedCount
## Min. : 0      Min. : 0      Min. : 0      Min. : 0
## 1st Qu.: 2524      1st Qu.: 258      1st Qu.: 277      1st Qu.: 4
## Median : 13042      Median : 661      Median : 915      Median : 22
## Mean : 45605      Mean : 1683      Mean : 10366      Mean : 143
## 3rd Qu.: 41709      3rd Qu.: 1678      3rd Qu.: 2296      3rd Qu.: 104
## Max. : 1742257      Max. : 68232      Max. : 6604309      Max. : 39981
## NA s : 29      NA s : 29      NA s : 29      NA s : 29
## acct_age      langDiv      mean_time_betwn_tweets      mCount
## Min. : 1361      Min. : 0.39      Min. : 0      Min. : 1.00
```

```
## 1st Qu.: 766986 1st Qu.: 2.95 1st Qu.: 40 1st Qu.: 1.00
## Median : 2275666 Median : 5.57 Median : 127 Median : 1.00
## Mean : 2221613 Mean : 5.59 Mean : 12156 Mean : 1.29
## 3rd Qu.: 3534304 3rd Qu.: 7.50 3rd Qu.: 549 3rd Qu.: 1.00
## Max. : 5209655 Max. : 29.07 Max. : 3598898 Max. : 4.00
## NA s : 193 NA s : 38
## bot
## 0: 2107
## 1: 511
##
##
##
##
##
```


Like before, we can evaluate many relationships simultaneously with `ggpairs`.

```
# inspect many trends with ggpairs
ggpairs(shredCVL)
```




Once we have some initial hypotheses we can make more specific plots.


```
ggplot(shredCVL, aes(x = followersCount, fill = bot)) +  
  geom_histogram()
```




```
# Some people have a lot of followers, but most don't. we need to lob off  
# the long tail so we can see the distribution better  
ggplot(filter(shredCVL, followersCount < 100),  
  aes(x = followersCount, fill = bot)) +  
  geom_histogram()
```




```
ggplot(filter(shredCVL, followersCount < 100),  
  aes(x = followersCount, fill = bot)) +  
  geom_histogram() +  
  facet_grid(bot ~.)
```




```
# how about the number of people they follow?
ggplot(shredCVL, aes(x = friendsCount, fill = bot)) +
  geom_histogram() +
  facet_grid(bot ~ .)
```




```
# it's a little hard to see
ggplot(filter(shredCVL, friendsCount < 2500),
  aes(x = friendsCount, fill = bot)) +
  geom_histogram() +
  facet_grid(bot ~.)
```

```
# what about account age?
ggplot(shredCVL, aes(x = acct_age, fill = bot)) +
  geom_histogram() +
  facet_grid(bot ~.)
```




```
# lexical diversity
ggplot(shredCVL, aes(x = langDiv, fill = bot)) +
  geom_histogram() +
  facet_grid(bot ~ .)
```



```
# what are the average values?
avg_diversity =
  shredCVL %>%
    group_by(bot) %>%
    summarize(avg_diversity = mean(langDiv, na.rm = TRUE))


# add it to the plot
ggplot(shredCVL, aes(x = langDiv, fill = bot)) +
  geom_histogram() +
  geom_vline(data = avg_diversity, aes(xintercept = avg_diversity)) +
  facet_grid(bot ~.)
```



Feature engineering

Feature engineering is the process of creating predictor variables using domain knowledge. We can test hypotheses about the importance of various relationships by creating new predictors that help interrogate those relationships. For example, you might hypothesize a relationship between the number of tweets made and the lexical diversity that is relevant to model. To test that, make a new categorical variable indicating whether an account holder is a 'heavy tweeter', 'medium tweeter' or 'light tweeter':

```
# the number of tweets per account has a long tail  
ggplot(shredCVL, aes(x = statusesCount)) +  
  geom_histogram()
```



```
# break into three categories by quantile
#quantile(shredCVL$statusesCount)


# low tweeters will be the bottom 25%,
shredCVL$tweet_volume = NA
shredCVL$tweet_volume = ifelse(shredCVL$statusesCount <= 188,
                                Light Tweeter ,
                                shredCVL$tweet_volume)

shredCVL$tweet_volume = ifelse((shredCVL$statusesCount > 188 & shredCVL$statusesCount <= 2646),
                                Medium Tweeter ,
                                shredCVL$tweet_volume)

shredCVL$tweet_volume = ifelse(shredCVL$statusesCount > 2646,
                                Heavy Tweeter ,
                                shredCVL$tweet_volume)

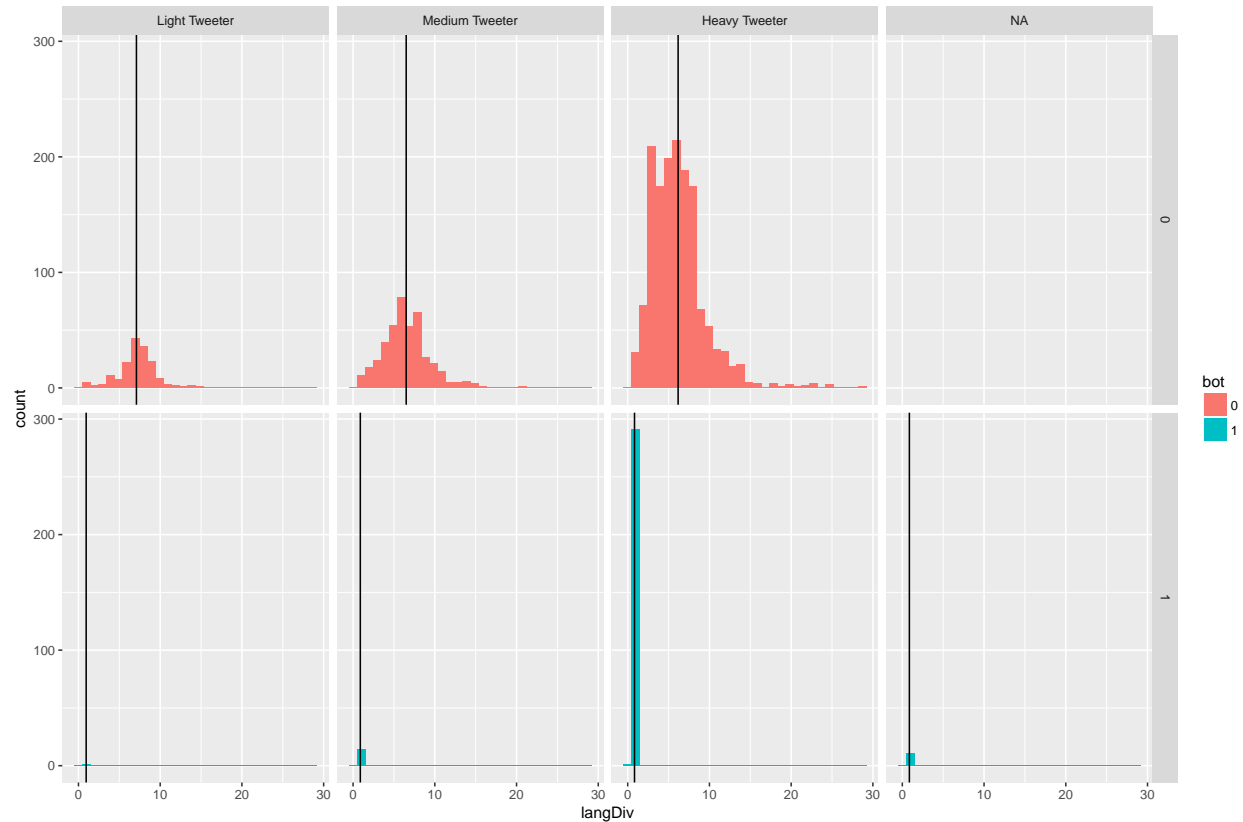
shredCVL$tweet_volume = factor(shredCVL$tweet_volume, levels = c( Light Tweeter , Medium Tweeter , Heavy Tweeter ))

# plot it!
ggplot(shredCVL, aes(x = statusesCount)) +
  geom_histogram(aes(fill = tweet_volume), bins = 100)
```



```
# update the figure
avg_diversity =
  shredCVL %>%
    group_by(bot, tweet_volume) %>%
    summarize(avg_diversity = mean(langDiv, na.rm = TRUE))

ggplot(shredCVL, aes(x = langDiv, fill = bot)) +
  geom_histogram() +
  geom_vline(data = avg_diversity, aes(xintercept = avg_diversity)) +
  facet_grid(bot ~ tweet_volume)
```



Logistic Regression

Training and testing sets

```
set.seed(243)
shredCVL = na.omit(shredCVL)

# select the training observations
in_train = createDataPartition(y = shredCVL$bot,
                                p = 0.75, # 75% in train, 25% in test
                                list = FALSE)

train = shredCVL[in_train, ]
test = shredCVL[-in_train, ]
```

Training logistic regressions


Check out this page for more types of logistic regression to try out.

```
logistic_model = train(bot ~ .,
                        data = train,
                        method = glm,
                        family = binomial,
                        preprocess = c(center, scale))
```

```
summary(logistic_model)
```

```
##
## Call:
## NULL
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.30      0.00      0.00      0.00      1.45
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -1.49e+02   5.18e+01  -2.87   0.0041
## statusesCount    8.40e+00   3.70e+00   2.27   0.0230
## friendsCount   -5.47e-01   5.87e-01  -0.93   0.3508
## followersCount  3.41e+01   1.64e+01   2.08   0.0378
## listedCount   -3.61e+01   2.19e+01  -1.65   0.0990
## acct_age       3.64e-03   1.07e+00   0.00   0.9973
## langDiv       -1.13e+02   3.92e+01  -2.89   0.0038
## mean_time_betwn_tweets 2.91e+00   5.76e+00   0.51   0.6133
## mCount         4.45e+00   3.73e+00   1.19   0.2323
## tweet_volumeMedium Tweeter 2.86e+00   2.24e+00   1.27   0.2026
## tweet_volumeHeavy Tweeter 9.31e-01   8.28e-01   1.12   0.2607
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1382.827  on 1811  degrees of freedom
## Residual deviance:   16.257  on 1801  degrees of freedom
## AIC: 38.26
##
## Number of Fisher Scoring iterations: 20
```

```
plot(varImp(logistic_model))
```

```
# test predictions
```

```
logistic_predictions = predict(logistic_model, newdata = test)
confusionMatrix(logistic_predictions, test$bot)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  0    1
```

```
##           0 523   0
```

```
##           1   3  76
```

```
##
```

```
##           Accuracy : 0.995
```

```
##           95% CI : (0.986, 0.999)
```

```
##           No Information Rate : 0.874
```

```
##           P-Value [Acc > NIR] : <2e-16
```

```
##
```

```
##           Kappa : 0.978
```

```
## McNemar's Test P-Value : 0.248
```

```
##
```

```
##           Sensitivity : 0.994
```

```
##           Specificity : 1.000
```

```
##           Pos Pred Value : 1.000
```

```
##           Neg Pred Value : 0.962
```

```
##           Prevalence : 0.874
```

```
##           Detection Rate : 0.869
```

```
##           Detection Prevalence : 0.869
```

```
##           Balanced Accuracy : 0.997
```

```
##
##      Positive Class : 0
##
```

There are subset selection methods for logistic regression as well. Try out method = glmStepAIC :

```
# stepwise logisitic regression
step_model = train(bot ~ .,
                   data = train,
                   method = glmStepAIC ,
                   family = binomial ,
                   preProcess = c( center , scale ))
```

```
summary(step_model)
```

```
##
## Call:
## NULL
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.16      0.00      0.00      0.00      1.44
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -149.93      46.73   -3.21  0.0013
## statusesCount     8.55       3.08    2.77  0.0056
## followersCount    17.12      13.25    1.29  0.1962
## listedCount     -37.52      16.91   -2.22  0.0265
## langDiv        -113.12      34.95   -3.24  0.0012
## mCount           4.54       3.49    1.30  0.1925
## tweet_volumeMedium Tweeter    2.19      1.98    1.10  0.2696
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1382.827  on 1811  degrees of freedom
## Residual deviance:   18.014  on 1805  degrees of freedom
## AIC: 32.01
##
## Number of Fisher Scoring iterations: 16
```

```
step_predictions = predict(step_model, newdata = test)
confusionMatrix(step_predictions, test$bot)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 523    0
##              1   3  76
##
##              Accuracy : 0.995
##              95% CI   : (0.986, 0.999)
##              No Information Rate : 0.874
##              P-Value [Acc > NIR] : <2e-16
##
```

```
##           Kappa : 0.978
## McNemar's Test P-Value : 0.248
##
##           Sensitivity : 0.994
##           Specificity : 1.000
##           Pos Pred Value : 1.000
##           Neg Pred Value : 0.962
##           Prevalence : 0.874
##           Detection Rate : 0.869
##           Detection Prevalence : 0.869
##           Balanced Accuracy : 0.997
##
##           Positive Class : 0
##
```

How do the models compare?

```
# compare
results = resamples(list(logistic_model = logistic_model,
                        step_model = step_model))
```

```
# compare accuracy and kappa
summary(results)
```

```
##
## Call:
## summary.resamples(object = results)
##
## Models: logistic_model, step_model
## Number of resamples: 25
##
## Accuracy
##           Min. 1st Qu. Median   Mean 3rd Qu.   Max. NA's
## logistic_model 0.9719 0.9869 0.9911 0.9896 0.9955 0.9985    0
## step_model     0.9881 0.9940 0.9956 0.9953 0.9970 1.0000    0
##
## Kappa
##           Min. 1st Qu. Median   Mean 3rd Qu.   Max. NA's
## logistic_model 0.8631 0.9458 0.9604 0.9523 0.9802 0.9934    0
## step_model     0.9496 0.9719 0.9803 0.9789 0.9872 1.0000    0
```

```
# plot results
dotplot(results)
```

